

Collaborative Metric Learning

Cheng-Kang Hsieh, Longqi Yang, Yin Cui, Tsung-Yi Lin, Serge Belongie, Deborah Estrin

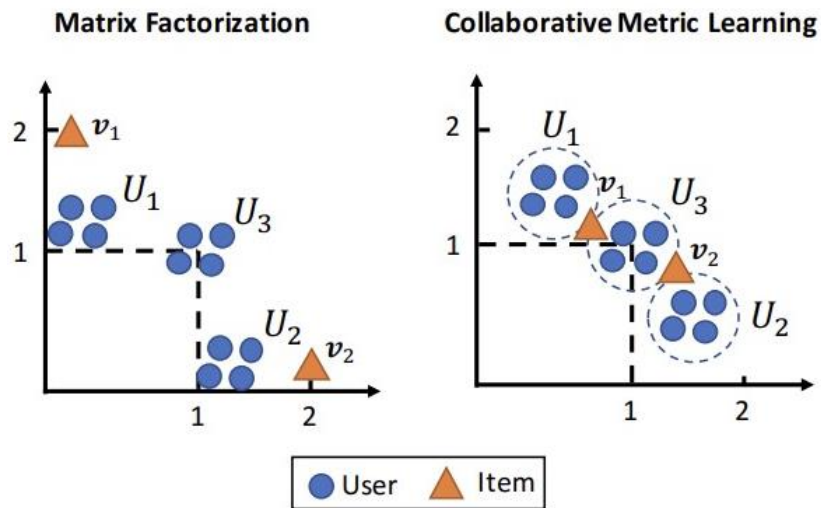
1. Introduction
2. BackGround
3. Model description
4. Evaluation
5. Conclusion

Introduction

- Distance is at the heart of many fundamental machine learning algorithms
 - Ex) K-nearest neighbor, K-means, SVMs
- Metric learning algorithms produce distance metric
 - similar & dissimilar
 - Assign smaller distances between similar pairs, larger distances between dissimilar pairs.

Introduction

- Triangle inequality
 - x is similar to y,z -> x pull two pairs closer, also pull the remaining pair (y,z) relatively close
 - matrix factorization -> does not satisfy the triangle inequality (dot product)




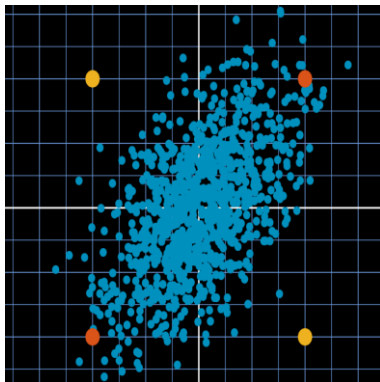
$$Matrix = \begin{bmatrix} 0 & 2 \\ 2 & 0 \\ 2 & 2 \end{bmatrix} \quad LatentUser = \begin{bmatrix} 0 & 1 \\ 1 & 0 \\ 1 & 1 \end{bmatrix} \quad LatentItem = \begin{bmatrix} 0 & 2 \\ 2 & 0 \end{bmatrix}$$

$$N \times M = (N \times K) \cdot (K \times M)$$

- what we want to know
 - users' preference
 - user-user & item-item similarity
 - > Collaborative Metric Learning
- explicit -> implicit

Background

- Metric Learning
 - $S = \{(x_i, x_j) \mid x_i \text{ and } x_j \text{ are considered similar}\}$
 - $D = \{(x_i, x_j) \mid x_i \text{ and } x_j \text{ are considered dissimilar}\}$
- Mahalanobis distance metric
 - $d_A(x_i, x_j) = \sqrt{(x_i - x_j)^T A (x_i - x_j)}$
 - where $A \in \mathbb{R}^{m \times m}$ is a positive semi-definite matrix



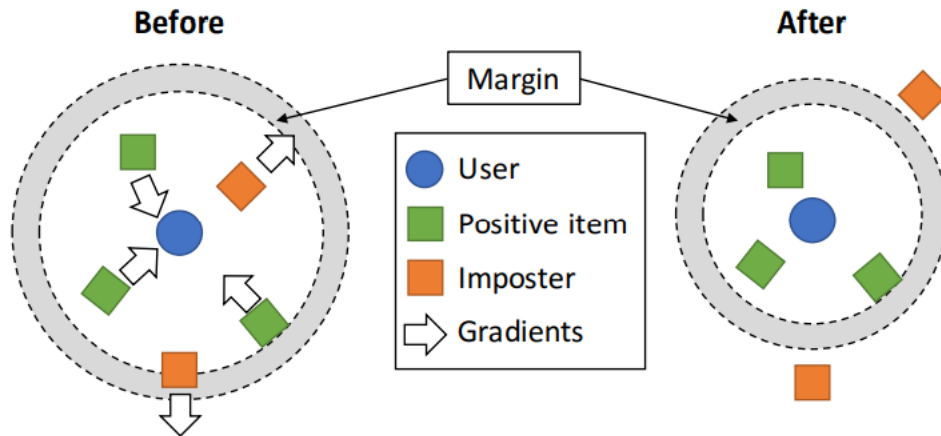
Background

- Metric Learning for kNN

- LMNN

- $\mathcal{L}_{pull}(d) = \sum_{j \sim i} d(x_i, x_j)^2$

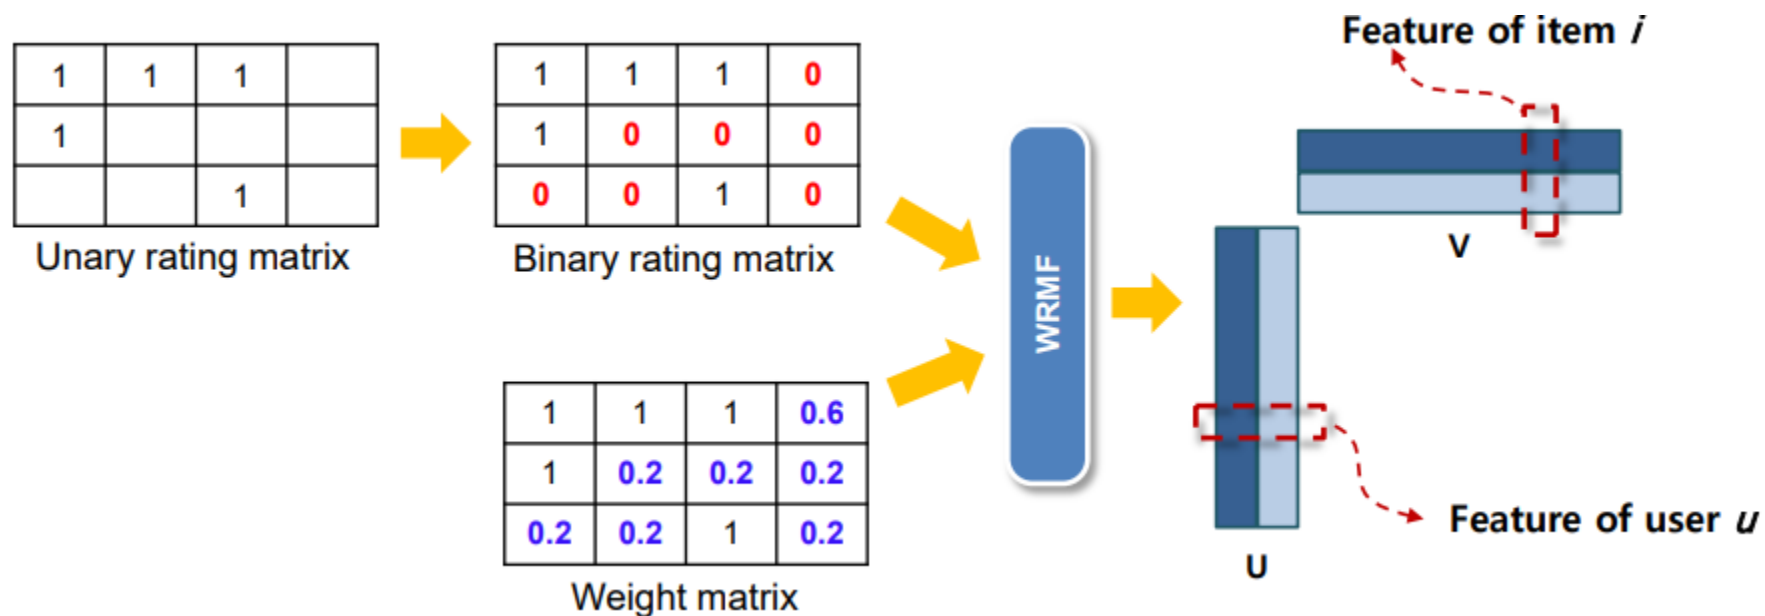
- $\mathcal{L}_{push}(d) = \sum_{j \sim i} \sum_k (1 - y_{ik}) \left[1 + d(x_i, x_j)^2 - d(x_i, x_k)^2 \right]_+$



Background

- Weighted regularized matrix factorization

$$\min_{\mathbf{u}_*, \mathbf{v}_*} \sum_{r_{ij} \in \mathcal{K}} c_{ij} (r_{ij} - \mathbf{u}_i^T \mathbf{v}_j)^2 + \lambda_u \|\mathbf{u}_i\|^2 + \lambda_v \|\mathbf{v}_j\|^2,$$



- Includes all the unobserved user-item interactions as negative samples
- Uses a case weight c_{ij} to reduce the impact of these uncertain samples

Background

- Bayesian Personalized Ranking

$$\min_{\mathbf{u}_*, \mathbf{v}_*} \sum_{i \in \mathcal{I}} \sum_{(j,k) \in \mathcal{D}_i} -\log \sigma(\mathbf{u}_i^T \mathbf{v}_j - \mathbf{u}_i^T \mathbf{v}_k) + \lambda_u \|\mathbf{u}_i\|^2 + \lambda_v \|\mathbf{v}_j\|^2$$

- (user, positive item, negative item)

Model description

- Model feature
 - Capture users' relative preferences for different items
 - Uses implicit data
 - Pulls the pairs in S closer and pushes the other pairs relatively further apart
 - By triangular inequality, will also cluster
 - User who co-like the same items together
 - Items that are co-liked by same users together

Model description

- Model formulation

$$\min_{\theta, \mathbf{u}_*, \mathbf{v}_*} \mathcal{L}_m + \lambda_f \mathcal{L}_f + \lambda_c \mathcal{L}_c$$

$$s. t. \|\mathbf{u}_*\|^2 \leq 1 \text{ and } \|\mathbf{v}_*\|^2 \leq 1$$

\mathcal{L}_m : Embedding Loss

\mathcal{L}_f : Feature Loss

\mathcal{L}_c : Covariance Loss

Model description

- Model formulation-Embedding Loss

$$\mathcal{L}_m(d) = \sum_{(i,j) \in S} \sum_{(i,k) \notin S} w_{ij} [m + D(i,j)^2 - d(i,k)^2]_+$$

- $m + d(i,j)^2 > d(i,k)^2 \rightarrow \text{loss}$
 - m : margin
 - w_{ij} : weight

Model description

- Model formulation-Embedding Loss

- Weighted Approximate-Rank Pairwise (WARP)
 - Penalize items at a lower rank

$$w_{ij} = \log(\text{rank}_d(i, j) + 1)$$

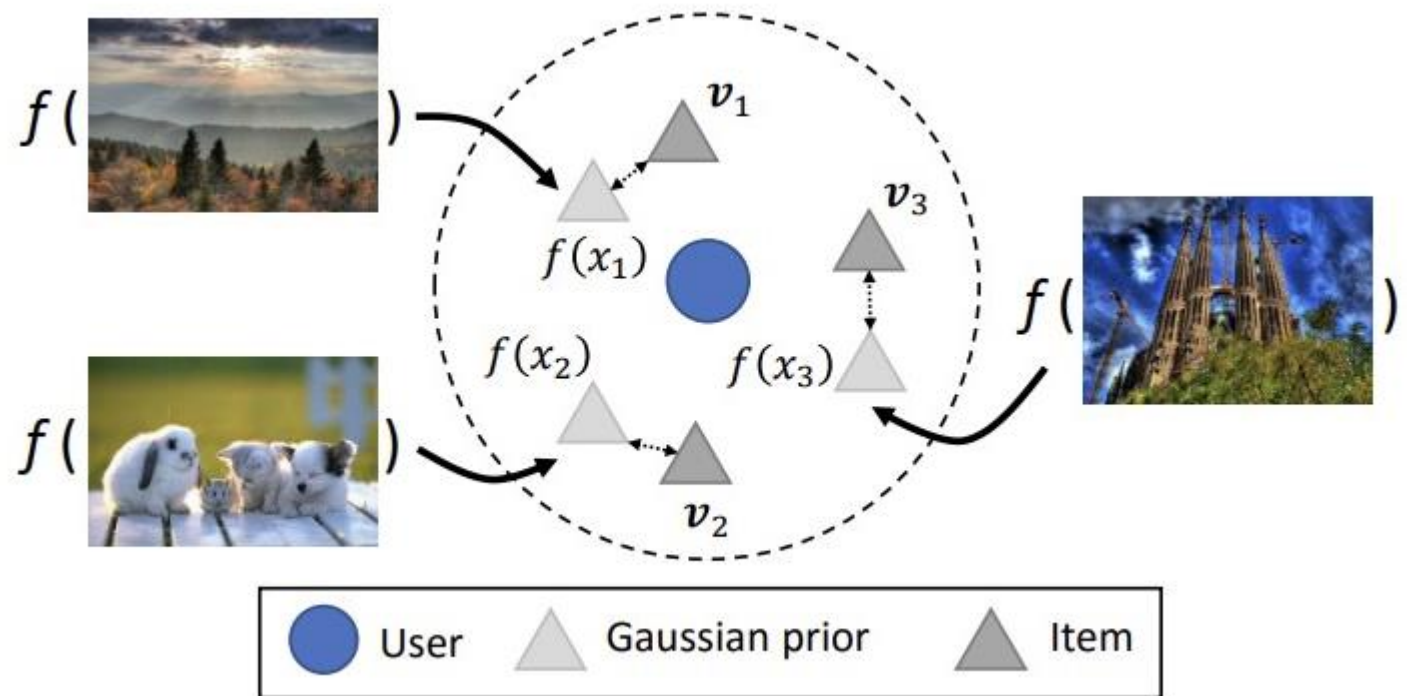
- Sample U negative items in parallel and compute the hinge loss
- Let M denote the number of impostors in U sample, $\text{rank}_d(i, j)$ is approximated to $\left\lfloor \frac{J \times M}{U} \right\rfloor$

Model description

- Model formulation-feature loss

$$\mathcal{L}_f(\theta, \mathbf{v}_*) = \sum_j ||f(\mathbf{x}_j, \theta) - \mathbf{v}_j||^2$$

- Think as transformation of input
- x_j : raw item j vector
- v_j : embedding item j vector
- $f(x)$: MLP



Model description

- Model formulation-regularization
 - kNN based model is known to be ineffective in a high-dimensional space if the data points spread too widely
 - > bound all the user/item
- Covariance regularization

$$\mathcal{L}_c = \frac{1}{N} (\|C\|_f^2 - \|\text{diag}(C)\|_2^2)$$

$$C_{ij} = \frac{1}{N} \sum_n (y_i^n - \mu_i)(y_j^n - \mu_j)$$

Model Evaluation

- Dataset

Table 1: Dataset Statistics.

	CiteULike	BookCX	Flickr	Medium	MovieLens20M	EchoNest
Domain	Paper	Book	Photography	News	Movie	Song
# Users	7,947	22,816	43,758	61,909	129,797	766,882
# Items	25,975	43,765	100,000	80,234	20,709	260,417
# Ratings	142,794	623,405	1,372,621	2,047,908	9,939,873	7,261,443
Concentration ^a	33.47%	33.10%	13.48%	55.38%	72.52%	65.88%
Features Type	Tags	Subjects	Image Features	Tags	Genres, Keywords	NA
# Feature Dim.	10,399	7,923	2,048	2,313	10,399	NA

Model Evaluation



- Evaluation

Table 2: Recall@50 and Recall@100 on the test set. (# dimensions $r = 100$) The best performing method is boldfaced. *, **, *** indicate $p \leq 0.05$, $p \leq 0.01$, and $p \leq 0.001$ based on the Wilcoxon signed rank test suggested in [41].

	WRMF	BPR	WARP	CML	<i>ours vs. best</i>	FM	VBPR	CDL	CML+F	<i>ours vs. best</i>
<i>Recall@50</i>										
CiteULike	0.2437	0.2489	0.1916	0.2714***	9.03%	0.1668	0.2807	0.3375**	0.3312	-1.86%
BookCX	0.0910	0.0812	0.0801	0.1037***	13.95%	0.1016	0.1004	0.0984	0.1147***	12.89%
Flickr	0.0667	0.0496	0.0576	0.0711***	6.59%	NA	0.0612	0.0679	0.0753***	10.89%
Medium	0.1457	0.1407	0.1619	0.1730***	6.41%	0.1298	0.1656	0.1682	0.1780***	5.82%
MovieLens	0.4317	0.3236	0.4649	0.4665	0.34%	0.4384	0.4521	0.4573	0.4617*	0.96%
EchoNest	0.2285	0.1246	0.2433	0.2460	1.10%	NA	NA	NA	NA	NA
<i>Recall@100</i>										
CiteULike	0.3112	0.3296	0.2526	0.3411***	3.37%	0.2166	0.3437	0.4173	0.4255**	1.96%
BookCX	0.1286	0.1230	0.1227	0.1436***	11.66%	0.1440	0.1455	0.1428	0.1712***	17.66%
Flickr	0.0821	0.0790	0.0797	0.0922***	12.30%	NA	0.0880	0.0909	0.1048***	15.29%
Medium	0.2112	0.2078	0.2336	0.2480***	6.16%	0.1900	0.2349	0.2408	0.2531***	5.10%
MovieLens	0.5649	0.4455	0.5989	0.6022	0.55%	0.5561	0.5712	0.5943	0.5976	0.55%
EchoNest	0.2891	0.1655	0.3021	0.3022	0.00%	NA	NA	NA	NA	NA

Model Evaluation

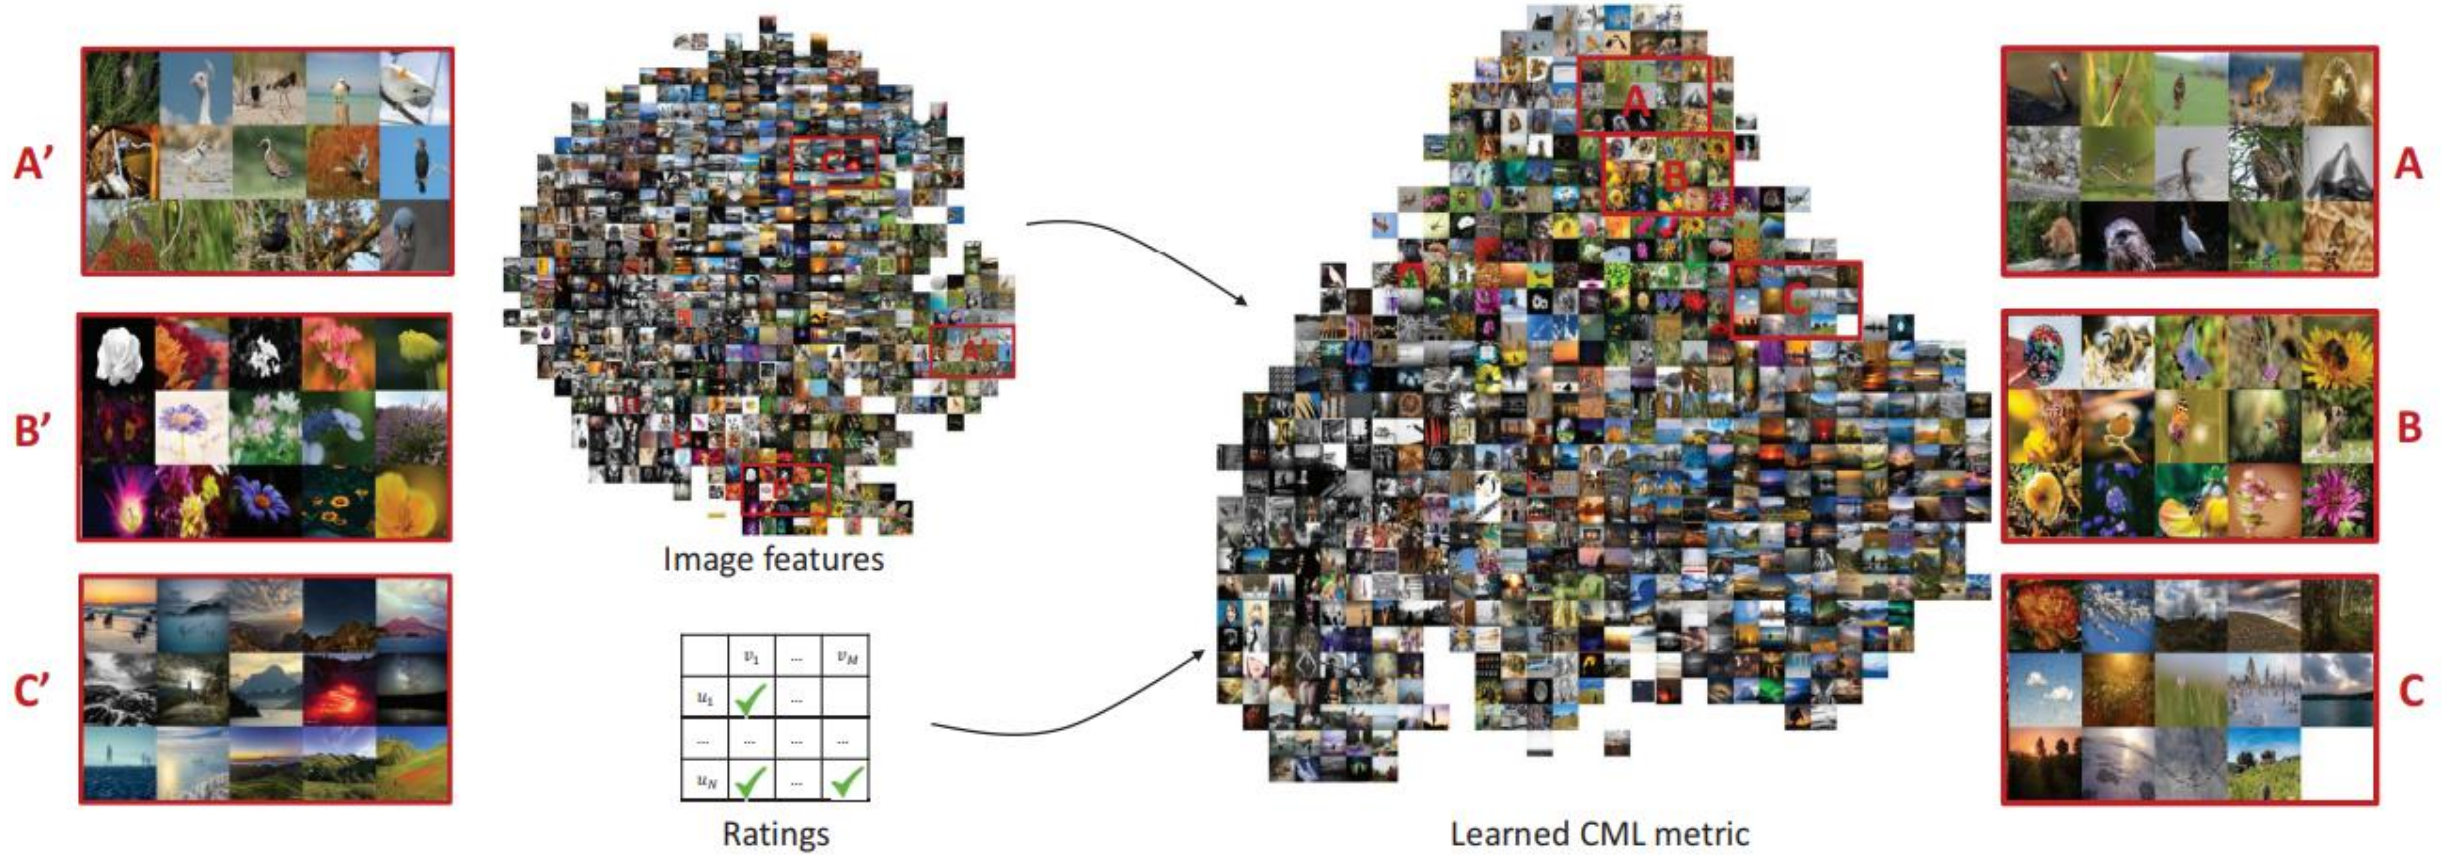
- Evaluation



- CML's throughput is improved by 106x with only 2% reduction in accuracy
- Over 8x faster than (optimized) MF models given the same accuracy

Model Evaluation

- Evaluation



Conclusion

- Conclusion

Task	Dataset	Model	Metric Name	Metric Value	Global Rank	Benchmark
Recommendation Systems	Million Song Dataset	CML	Recall@50	0.2460	# 6	Compare
			Recall@100	0.3022	# 1	Compare
Recommendation Systems	MovieLens 1M	CML	HR@10	0.7216	# 6	Compare
			nDCG@10	0.5413	# 5	Compare
Recommendation Systems	MovieLens 20M	CML	Recall@50	0.4665	# 9	Compare
			HR@10	0.7764	# 3	Compare
			nDCG@10	0.5301	# 3	Compare
Recommendation Systems	Netflix	CML	nDCG@10	0.2948	# 3	Compare
			Recall@10	0.4612	# 2	Compare