

# GraphSAGE

## Inductive Representation Learning on Large Graphs

William L. Hamilton, Rex Ying and Jure Leskovec

---

**NIPS 2017**  
Paper Review

Jihwan Oh

2023.02.07

# Contents

---

1. Introduction
2. Related Works
3. Model: GraphSAGE
4. Experiment
5. Conclusion
6. Implement

## 1. Limitation of previous tasks

Low-dimensional vector embeddings is useful for graph prediction and analysis.

**However**, it has focused on embedding nodes from a **single fixed graph (Transductive)**

Real-world applications require embeddings to be quickly generated for unseen nodes,  
or entirely new (sub)graphs

**Therefore**, we should consider the **generalization** to unseen node. (**Inductive**)

# Introduction

2

## 2. Transductive vs Inductive Learning






Figure 1. Node classification in transductive setting. At training time, the learning algorithm has access to all the nodes and edges including nodes for which labels are to be predicted.

### Transductive Learning



(a) A model  $\mathcal{H}$  is learned over some graph



(b) The model is then by applied to new nodes and edges

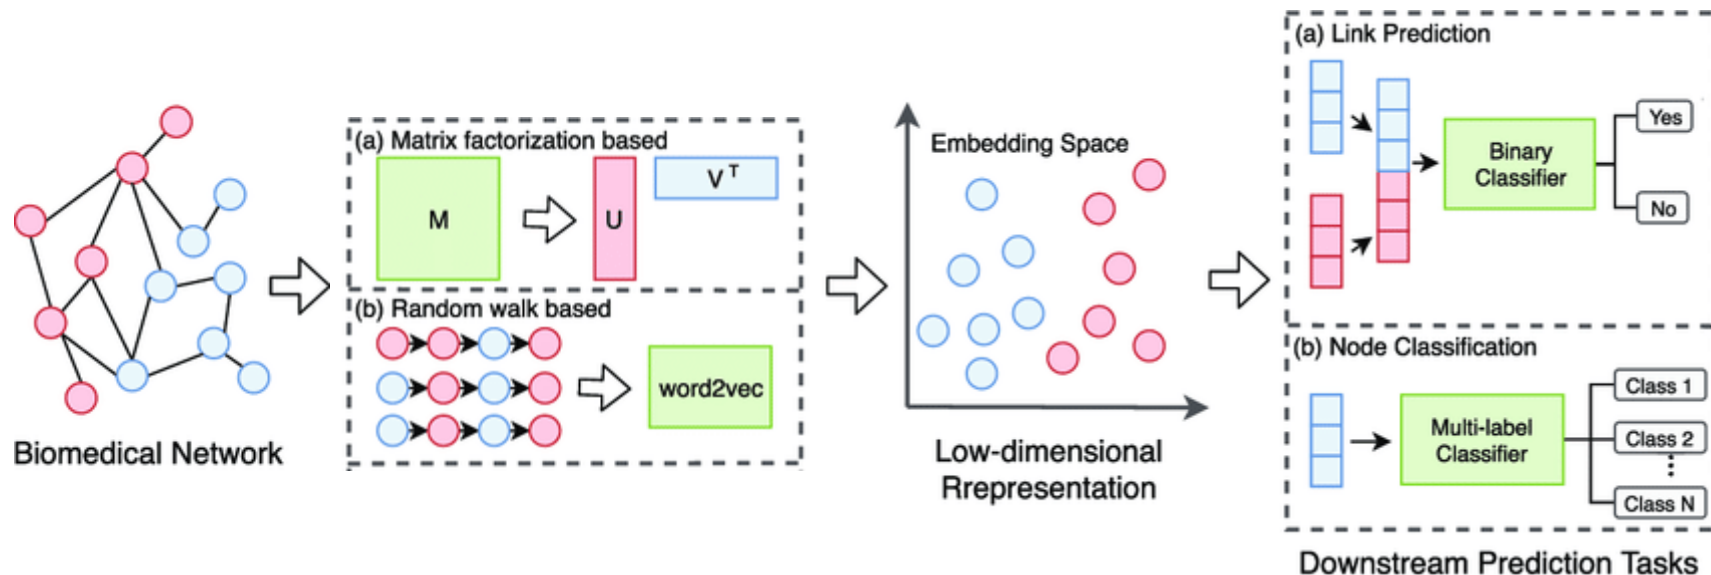
Figure 2. Node classification in inductive setting. Once learned, the model can be applied to new unseen nodes (outlined in red). There may or may not exist edges between such new nodes and the nodes used for training.

### Inductive Learning

## 1. Factorization-based embedding approaches

Low-dimensional embeddings using random walk statistics and matrix factorization-based learning objectives


- Directly train node embeddings for individual nodes (Transductive)
- Require expensive additional training to make predictions on new nodes



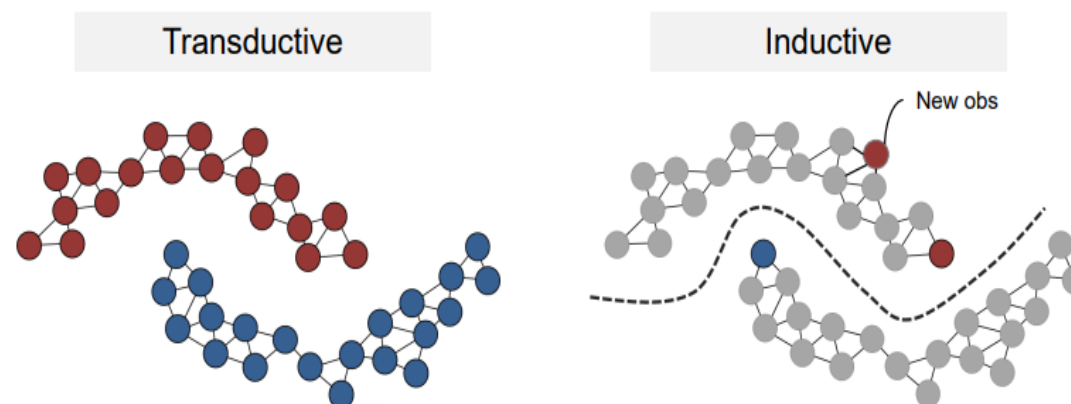
## 2. Graph convolutional networks (GCN)

Graph convolutional networks (GCNs) have only been applied in the transductive setting with fixed graphs

**GraphSAGE** can be viewed as an extension of the **GCN framework to the inductive setting**



(a) Graph Convolutional Network



## 1. Proposed model: GraphSAGE

GraphSAGE: SAmple and aggreGatE

- Generalized embedding (Inductive)
- Using **aggregate function** from neighbor nodes
- Both applied to supervised and unsupervised learning
- Both learn distribution and topological structure in neighbor nodes.
- Low computational cost
- Better performance than previous tasks (DeepWalk, GCN, etc.)

## 2. Embedding generation algorithm

Assume: Model already trained = Parameter fixed

Parameter =  $\text{AGGREGATE}_K$ : aggregator function,  $\mathbf{W}^k$ : weight matrix

---

**Algorithm 1:** GraphSAGE embedding generation (i.e., forward propagation) algorithm

---

**Input** : Graph  $\mathcal{G}(\mathcal{V}, \mathcal{E})$ ; input features  $\{\mathbf{x}_v, \forall v \in \mathcal{V}\}$ ; depth  $K$ ; weight matrices  $\mathbf{W}^k, \forall k \in \{1, \dots, K\}$ ; non-linearity  $\sigma$ ; differentiable aggregator functions  $\text{AGGREGATE}_k, \forall k \in \{1, \dots, K\}$ ; neighborhood function  $\mathcal{N} : v \rightarrow 2^{\mathcal{V}}$

**Output** : Vector representations  $\mathbf{z}_v$  for all  $v \in \mathcal{V}$

```
1  $\mathbf{h}_v^0 \leftarrow \mathbf{x}_v, \forall v \in \mathcal{V}$  ;  
2 for  $k = 1 \dots K$  do  
3   for  $v \in \mathcal{V}$  do  
4      $\mathbf{h}_{\mathcal{N}(v)}^k \leftarrow \text{AGGREGATE}_k(\{\mathbf{h}_u^{k-1}, \forall u \in \mathcal{N}(v)\})$ ;  
5      $\mathbf{h}_v^k \leftarrow \sigma(\mathbf{W}^k \cdot \text{CONCAT}(\mathbf{h}_v^{k-1}, \mathbf{h}_{\mathcal{N}(v)}^k))$   
6   end  
7    $\mathbf{h}_v^k \leftarrow \mathbf{h}_v^k / \|\mathbf{h}_v^k\|_2, \forall v \in \mathcal{V}$   
8 end  
9  $\mathbf{z}_v \leftarrow \mathbf{h}_v^K, \forall v \in \mathcal{V}$ 
```

---



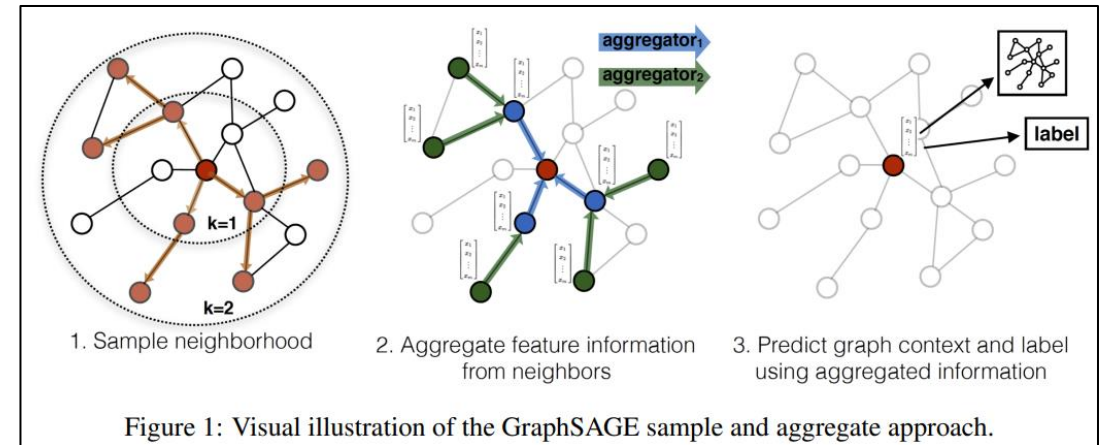
## 2. Embedding generation algorithm

**Algorithm 1:** GraphSAGE embedding generation (i.e., forward propagation) algorithm

**Input** : Graph  $\mathcal{G}(\mathcal{V}, \mathcal{E})$ ; input features  $\{\mathbf{x}_v, \forall v \in \mathcal{V}\}$ ; depth  $K$ ; weight matrices  $\mathbf{W}^k, \forall k \in \{1, \dots, K\}$ ; non-linearity  $\sigma$ ; differentiable aggregator functions  $\text{AGGREGATE}_k, \forall k \in \{1, \dots, K\}$ ; neighborhood function  $\mathcal{N} : v \rightarrow 2^{\mathcal{V}}$

**Output** : Vector representations  $\mathbf{z}_v$  for all  $v \in \mathcal{V}$

```
1  $\mathbf{h}_v^0 \leftarrow \mathbf{x}_v, \forall v \in \mathcal{V}$  ;  
2 for  $k = 1 \dots K$  do  
3   for  $v \in \mathcal{V}$  do  
4      $\mathbf{h}_{\mathcal{N}(v)}^k \leftarrow \text{AGGREGATE}_k(\{\mathbf{h}_u^{k-1}, \forall u \in \mathcal{N}(v)\})$ ;  
5      $\mathbf{h}_v^k \leftarrow \sigma(\mathbf{W}^k \cdot \text{CONCAT}(\mathbf{h}_v^{k-1}, \mathbf{h}_{\mathcal{N}(v)}^k))$   
6   end  
7    $\mathbf{h}_v^k \leftarrow \mathbf{h}_v^k / \|\mathbf{h}_v^k\|_2, \forall v \in \mathcal{V}$   
8 end  
9  $\mathbf{z}_v \leftarrow \mathbf{h}_v^K, \forall v \in \mathcal{V}$ 
```



## 2. Embedding generation algorithm (mini batch)

**Algorithm 2:** GraphSAGE minibatch forward propagation algorithm

**Input** : Graph  $\mathcal{G}(\mathcal{V}, \mathcal{E})$ ;  
input features  $\{\mathbf{x}_v, \forall v \in \mathcal{B}\}$ ;  
depth  $K$ ; weight matrices  $\mathbf{W}^k, \forall k \in \{1, \dots, K\}$ ;  
non-linearity  $\sigma$ ;  
differentiable aggregator functions  $\text{AGGREGATE}_k, \forall k \in \{1, \dots, K\}$ ;  
neighborhood sampling functions,  $\mathcal{N}_k : v \rightarrow 2^{\mathcal{V}}, \forall k \in \{1, \dots, K\}$

**Output** : Vector representations  $\mathbf{z}_v$  for all  $v \in \mathcal{B}$

```
1  $\mathcal{B}^K \leftarrow \mathcal{B}$ ;  
2 for  $k = K \dots 1$  do  
3    $\mathcal{B}^{k-1} \leftarrow \mathcal{B}^k$ ;  
4   for  $u \in \mathcal{B}^k$  do  
5      $\mathcal{B}^{k-1} \leftarrow \mathcal{B}^{k-1} \cup \mathcal{N}_k(u)$ ;  
6   end  
7 end  
8  $\mathbf{h}_u^0 \leftarrow \mathbf{x}_v, \forall v \in \mathcal{B}^0$ ;  
9 for  $k = 1 \dots K$  do  
10  for  $u \in \mathcal{B}^k$  do  
11     $\mathbf{h}_{\mathcal{N}(u)}^k \leftarrow \text{AGGREGATE}_k(\{\mathbf{h}_{u'}^{k-1}, \forall u' \in \mathcal{N}_k(u)\})$ ;  
12     $\mathbf{h}_u^k \leftarrow \sigma(\mathbf{W}^k \cdot \text{CONCAT}(\mathbf{h}_u^{k-1}, \mathbf{h}_{\mathcal{N}(u)}^k))$ ;  
13     $\mathbf{h}_u^k \leftarrow \mathbf{h}_u^k / \|\mathbf{h}_u^k\|_2$ ;  
14  end  
15 end  
16  $\mathbf{z}_u \leftarrow \mathbf{h}_u^K, \forall u \in \mathcal{B}$ 
```

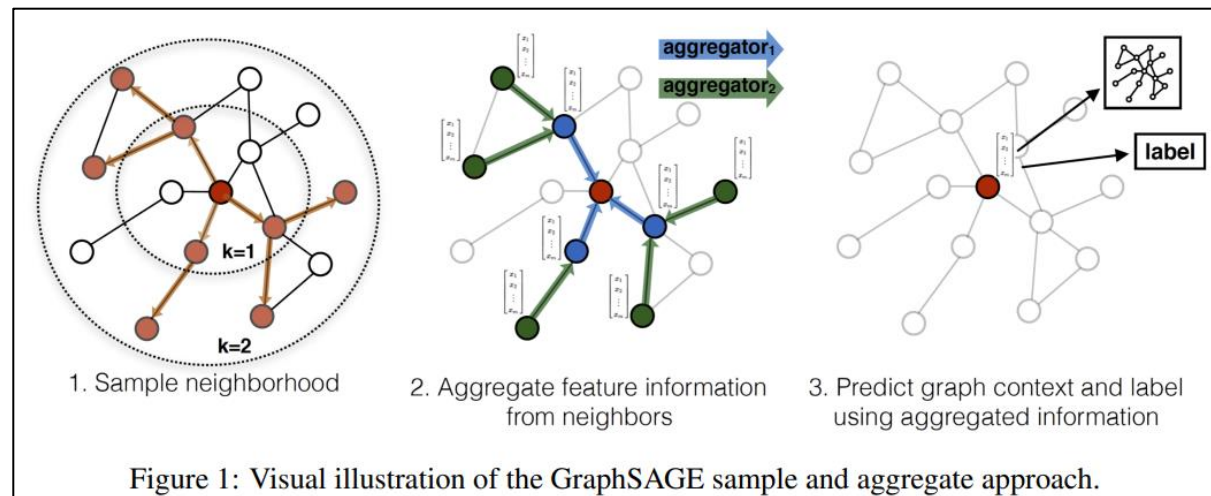
Make each mini batch  $\mathcal{B}^k, k \in \{1, \dots, K\}$

Identical to previous pseudocode

## 3. Neighborhood Definition

Use a fixed-size set of neighbors in order to keep computational cost.

- Without this sampling,  $O(V)$   $\rightarrow$  High computational cost
- With this sampling,  $O(\prod_{i=1}^k S_i)$ ,  $S_i$  is the size of neighborhood set for  $i \in \{1, \dots, K\}$
- $K=2$ ,  $S_1 \times S_2 \leq 500$  is working well.
- Low computational cost



## 4. Aggregator Architectures

In Graph, nodes and its neighbors have no ordering.

→ Aggregator should be symmetric. (Permutation invariant)

Symmetric, Trainable, High representational capacity

- Mean Aggregator
- LSTM Aggregator
- Pooling Aggregator

## 4. Aggregator Architectures



In Graph, nodes and its neighbors have no ordering.

→ Aggregator should be symmetric. (Permutation invariant)

Symmetric, Trainable, High representational capacity

- Mean Aggregator
- LSTM Aggregator
- Pooling Aggregator

```
1  $\mathbf{h}_v^0 \leftarrow \mathbf{x}_v, \forall v \in \mathcal{V}$ ;  
2 for  $k = 1 \dots K$  do  
3   for  $v \in \mathcal{V}$  do  
4      $\mathbf{h}_{\mathcal{N}(v)}^k \leftarrow \text{AGGREGATE}_k(\{\mathbf{h}_u^{k-1}, \forall u \in \mathcal{N}(v)\})$ ;  
5      $\mathbf{h}_v^k \leftarrow \sigma(\mathbf{W}^k \cdot \text{CONCAT}(\mathbf{h}_v^{k-1}, \mathbf{h}_{\mathcal{N}(v)}^k))$   
6   end  
7    $\mathbf{h}_v^k \leftarrow \mathbf{h}_v^k / \|\mathbf{h}_v^k\|_2, \forall v \in \mathcal{V}$   
8 end  
9  $\mathbf{z}_v \leftarrow \mathbf{h}_v^K, \forall v \in \mathcal{V}$ 
```



## 4. Aggregator Architectures


### (1) Mean aggregator

- Similar to GCN, but it use Concatenating instead of Adding.

```

1  $\mathbf{h}_v^0 \leftarrow \mathbf{x}_v, \forall v \in \mathcal{V}$ ;
2 for  $k = 1 \dots K$  do
3   for  $v \in \mathcal{V}$  do
4      $\mathbf{h}_{\mathcal{N}(v)}^k \leftarrow \text{AGGREGATE}_k(\{\mathbf{h}_u^{k-1}, \forall u \in \mathcal{N}(v)\})$ ;
5      $\mathbf{h}_v^k \leftarrow \sigma(\mathbf{W}^k \cdot \text{CONCAT}(\mathbf{h}_v^{k-1}, \mathbf{h}_{\mathcal{N}(v)}^k))$ 
6   end
7    $\mathbf{h}_v^k \leftarrow \mathbf{h}_v^k / \|\mathbf{h}_v^k\|_2, \forall v \in \mathcal{V}$ 
8 end
9  $\mathbf{z}_v \leftarrow \mathbf{h}_v^K, \forall v \in \mathcal{V}$ 
    
```

$$\text{AGGREGATE}_K : \mathbf{h}_v^k \leftarrow \sigma(\mathbf{W} \cdot \text{MEAN}(\{\mathbf{h}_v^{k-1}\} \cup \{\mathbf{h}_u^{k-1}, \forall u \in \mathcal{N}(v)\}))$$





## 4. Aggregator Architectures

### (2) LSTM aggregator

LSTM: Long Short-Term Memory

- Variant of RNN (Better performance than original RNN)
- Advantage of **larger expressive capability**
- It is not permutation invariant since their inputs are sequential.

→ Applying the LSTMs to a **random permutation** of the node's neighbors (**Permutation invariant**)




## 4. Aggregator Architectures

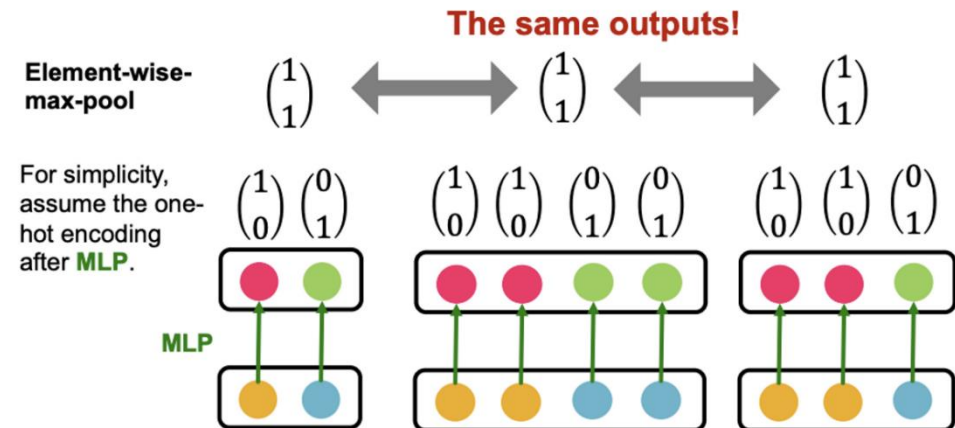
### (3) Pooling aggregator

- Each neighbor's vector is independently fed through MLP, then do element-wise max pooling.

$$\text{AGGREGATE}_k^{\text{pool}} = \max(\{\sigma(\mathbf{W}_{\text{pool}} \mathbf{h}_{u_i}^k + \mathbf{b}), \forall u_i \in \mathcal{N}(v)\})$$



GCN : Mean pooling



GraphSAGE : Max pooling



## 5. Learning the parameters of GraphSAGE

**Unsupervised Learning** : Graph-based loss

$$J_{\mathcal{G}}(\mathbf{z}_u) = -\log(\sigma(\mathbf{z}_u^{\top} \mathbf{z}_v)) - Q \cdot \mathbb{E}_{v_n \sim P_n(v)} \log(\sigma(-\mathbf{z}_u^{\top} \mathbf{z}_{v_n}))$$

- $v$  : Node that co-occur near  $u$  on fixed-length random walk
- $\mathbf{z}_u$  : Output representation by algorithm
- $P_n$  : Negative sampling distribution
- $v_n$  : Negative samples
- $Q$  : Number of negative samples

**Supervised Learning** : Cross-Entropy loss

## 5. Learning the parameters of GraphSAGE

**Unsupervised Learning** : Graph-based loss

$$J_{\mathcal{G}}(\mathbf{z}_u) = -\log(\sigma(\mathbf{z}_u^{\top} \mathbf{z}_v)) - Q \cdot \mathbb{E}_{v_n \sim P_n(v)} \log(\sigma(-\mathbf{z}_u^{\top} \mathbf{z}_{v_n}))$$

- $v$  : Node that co-occur near  $u$  on fixed-length random walk
- $\mathbf{z}_u$  : Output representation by algorithm
- $P_n$  : Negative sampling distribution
- $v_n$  : Negative samples
- $Q$  : Number of negative samples

**Supervised Learning** : Cross-Entropy loss

Train **Aggregator function** and **Weight matrix** with SGD

```
1  $\mathbf{h}_v^0 \leftarrow \mathbf{x}_v, \forall v \in \mathcal{V}$ ;  
2 for  $k = 1 \dots K$  do  
3   for  $v \in \mathcal{V}$  do  
4      $\mathbf{h}_{\mathcal{N}(v)}^k \leftarrow \text{AGGREGATE}_k(\{\mathbf{h}_u^{k-1}, \forall u \in \mathcal{N}(v)\})$ ;  
5      $\mathbf{h}_v^k \leftarrow \sigma(\mathbf{W}^k \cdot \text{CONCAT}(\mathbf{h}_v^{k-1}, \mathbf{h}_{\mathcal{N}(v)}^k))$   
6   end  
7    $\mathbf{h}_v^k \leftarrow \mathbf{h}_v^k / \|\mathbf{h}_v^k\|_2, \forall v \in \mathcal{V}$   
8 end  
9  $\mathbf{z}_v \leftarrow \mathbf{h}_v^K, \forall v \in \mathcal{V}$ 
```

## 1. Models and Loss function

### Models

- **GraphSAGE : GCN, Mean, LSTM, Max polling** ( $K=2, S_1 \times S_2 \leq 500$ )
- Random classifier
- Logistic Regression feature-based classifier (Ignore graph structure)
- DeepWalk
- DeepWalk + Raw features (Logistic Regression)

### Loss function

- Supervised learning: Cross-entropy loss
- Unsupervised learning: Graph-based loss

## 2. Dataset

### Citation Data

- Predicting paper subject categories on a large citation dataset
- Train : Test = 8:2 (approximate)

### Reddit Data

- Predict which community different Reddit posts belong to
- Train : Test = 8:2 (approximate)

### PPI (Protein-Protein Interaction) Data

- Predict protein-protein interactions
- For multi-graph generalization

# Experiment

19

## 3. Results

1. GraphSAGE performed better than benchmark tasks
2. Performance : LSTM, Max pooling > mean > GCN

\*Environment: Non-linear activation function: **ReLU**, Optimizer: **Adam**

Name	Citation		Reddit		PPI	
	Unsup. F1	Sup. F1	Unsup. F1	Sup. F1	Unsup. F1	Sup. F1
Random	0.206	0.206	0.043	0.042	0.396	0.396
Raw features	0.575	0.575	0.585	0.585	0.422	0.422
DeepWalk	0.565	0.565	0.324	0.324	—	—
DeepWalk + features	0.701	0.701	0.691	0.691	—	—
GraphSAGE-GCN	0.742	0.772	<b>0.908</b>	0.930	0.465	0.500
GraphSAGE-mean	0.778	0.820	0.897	0.950	0.486	0.598
GraphSAGE-LSTM	0.788	0.832	<b>0.907</b>	<b>0.954</b>	0.482	<b>0.612</b>
GraphSAGE-pool	<b>0.798</b>	<b>0.839</b>	0.892	0.948	<b>0.502</b>	0.600
% gain over feat.	39%	46%	55%	63%	19%	45%

Table 1: Prediction results for the three datasets (micro-averaged F1 scores). Results for unsupervised and fully supervised GraphSAGE are shown. Analogous trends hold for macro-averaged scores.

# Experiment

20

## 3. Results

**A:** GraphSAGE is faster than benchmark tasks when training data

**B:** Neighborhood sample size and accuracy ( $K=2$ ,  $S_1 = S_2$ )

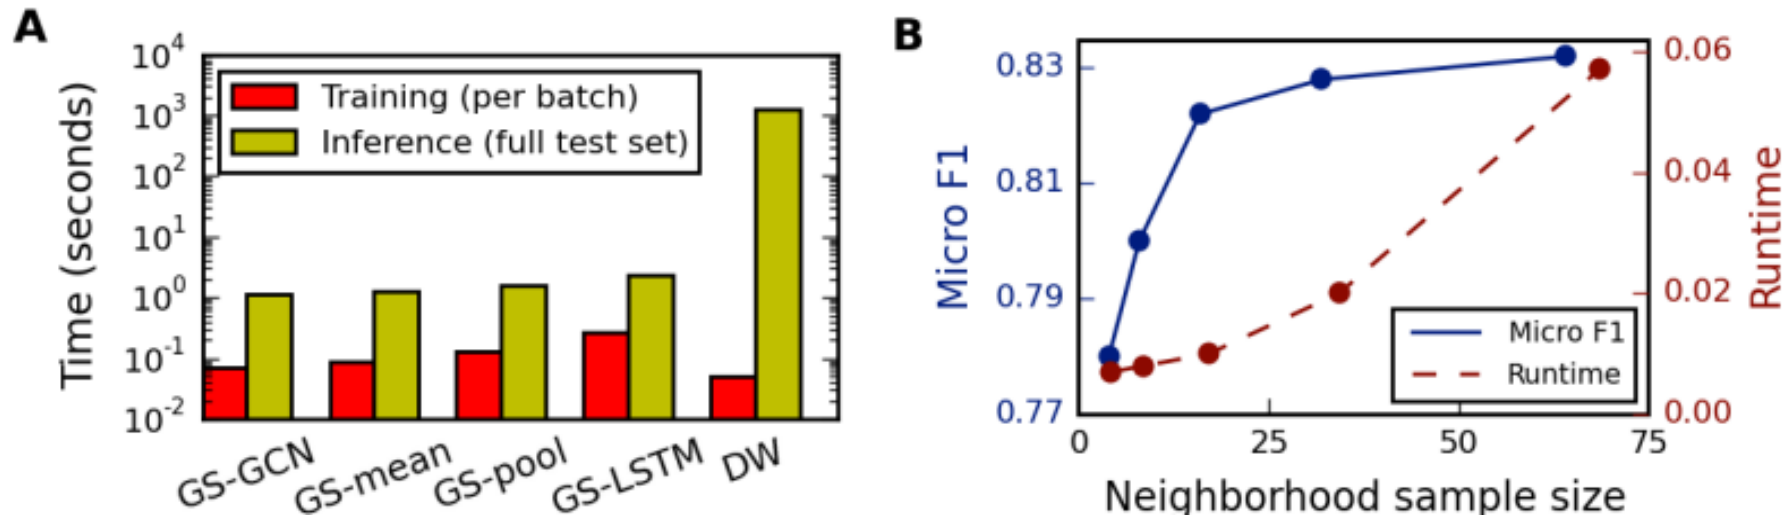


Figure 2: **A:** Timing experiments on Reddit data, with training batches of size 512 and inference on the full test set (79,534 nodes). **B:** Model performance with respect to the size of the sampled neighborhood, where the “neighborhood sample size” refers to the number of neighbors sampled at each depth for  $K = 2$  with  $S_1 = S_2$  (on the citation data using GraphSAGE-mean).

# Conclusion

## GraphSAGE

1. GraphSAGE is efficient algorithm for generating embeddings from unseen nodes.  
(**Inductive Learning**)
2. Effectively trade off performance and runtime in large graphs.
3. A number of extensions and potential improvements are possible,  
such as extending GraphSAGE to incorporate directed or multi-modal graphs.

# Implement

22

## Environment Setting

```
1 import torch
2 import dgl → Deep graph library
3 from dgl.data import CoraGraphDataset #Data 1
4 from dgl.data import RedditDataset #Data 2
5 #from dgl.data import PPIDataset #Data 3 (Not working)
6
7 from dgl.nn import SAGEConv → GraphSAGE layer
8
9 import matplotlib.pyplot as plt
10 import numpy as np
11
12 import networkx as nx
13 from torch.nn.parameter import Parameter
14 from torch.nn.modules.module import Module
15
16 import scipy
17 import scipy.sparse as sp
18 import torch.nn as nn
19 import torch.nn.functional as F
```

→ Data Loading

$$\begin{aligned}h_{\mathcal{N}(i)}^{(l+1)} &= \text{aggregate}(\{h_j^l, \forall j \in \mathcal{N}(i)\}) \\h_i^{(l+1)} &= \sigma(W \cdot \text{concat}(h_i^l, h_{\mathcal{N}(i)}^{l+1})) \\h_i^{(l+1)} &= \text{norm}(h_i^{(l+1)})\end{aligned}$$



## Environment Setting

```
✓ [6] 1 torch.__version__  
31 초  
'1.13.1+cu116'
```

→ Check pytorch version

```
✓ [1] 1 pip install --pre dgl-cu116 -f https://data.dgl.ai/wheels-test/repo.html  
31 초  
Looking in indexes: https://pypi.org/simple, https://us-python.pkg.dev/colab-wheels/public/simple/  
Looking in links: https://data.dgl.ai/wheels-test/repo.html  
Collecting dgl-cu116  
  Downloading https://data.dgl.ai/wheels-test/dgl_cu116-1.0a230116-cp38-cp38-manylinux1_x86_64.whl (265.6 MB)  
----- 265.6/265.6 MB 4.1 MB/s eta 0:00:00  
Requirement already satisfied: requests>=2.19.0 in /usr/local/lib/python3.8/dist-packages (from dgl-cu116) (2.25.1)  
Requirement already satisfied: numpy>=1.14.0 in /usr/local/lib/python3.8/dist-packages (from dgl-cu116) (1.21.6)  
Collecting psutil>=5.8.0  
  Downloading psutil-5.9.4-cp36-abi3-manylinux_2_12_x86_64.manylinux2010_x86_64.manylinux2014_x86_64.whl (280 kB)  
----- 280.2/280.2 KB 6.0 MB/s eta 0:00:00
```

→ Download dgl cuda

## Dataset

✓  
0초

```
[6] 1 dataset1 = CoraGraphDataset(verbose=True)
```

```
Downloading /root/.dgl/cora_v2.zip from https://data.dgl.ai/dataset/cora\_v2.zip...  
Extracting file to /root/.dgl/cora_v2  
Finished data loading and preprocessing.  
  NumNodes: 2708  
  NumEdges: 10556  
  NumFeats: 1433  
  NumClasses: 7  
  NumTrainingSamples: 140  
  NumValidationSamples: 500  
  NumTestSamples: 1000  
Done saving data into cached files.
```

**Cora Dataset**

✓  
1분

```
[7] 1 dataset2 = RedditDataset(self_loop=False, verbose=True)
```

```
Downloading /root/.dgl/reddit.zip from https://data.dgl.ai/dataset/reddit.zip...  
Extracting file to /root/.dgl/reddit  
Finished data loading.  
  NumNodes: 232965  
  NumEdges: 114615892  
  NumFeats: 602  
  NumClasses: 41  
  NumTrainingSamples: 153431  
  NumValidationSamples: 23831  
  NumTestSamples: 55703  
Done saving data into cached files.
```

**Reddit Dataset**


# Implement

25

## Model : GraphSAGE

```
1 class SAGE(nn.Module):
2     def __init__(self, in_size, hid_size, out_size):
3         super().__init__()
4         self.layers = nn.ModuleList()
5         # two-layer GraphSAGE-gcn, mean, pool, lstm
6         self.layers.append(SAGEConv(in_size, hid_size, "gcn"))
7         self.layers.append(SAGEConv(hid_size, out_size, "gcn"))
8         self.dropout = nn.Dropout(0.5)
9
10    def forward(self, graph, x):
11        h = x
12        for l, layer in enumerate(self.layers):
13            h = layer(graph, h)
14            if l != len(self.layers) - 1:
15                h = F.relu(h)
16                h = self.dropout(h)
17        return h
```

Two layers(K=2)



# Implement

26

## Evaluate function : F1-micro

```
19 # f1-micro
20 def evaluate(g, features, labels, mask, model):
21     model.eval()
22     with torch.no_grad():
23         y_actual = labels
24         y_pred = model(g, features)
25         y_pred = y_pred[mask]
26         y_actual = y_actual[mask]
27         _, indices = torch.max(y_pred, dim=1)
28         correct = torch.sum(indices == y_actual)
29         incorrect = torch.sum(indices != y_actual)
30         #f1-micro = TP/(TP*0.5(FP+FN))
31         return correct.item() / (correct.item() + 0.5 * incorrect.item())
```

# Implement

27

## Train (Full batch)

```
34 def train(g, features, labels, masks, model):
35     # define train/val samples, loss function and optimizer
36     train_mask, val_mask = masks
37     loss_fcn = nn.CrossEntropyLoss()
38     optimizer = torch.optim.Adam(model.parameters(), lr=1e-2, weight_decay=5e-4)
39
40     y_actual = labels
41     accuracy_list = []
42     loss_list = []
43     # training loop
44     for epoch in range(200):
45         model.train()
46         y_pred = model(g, features)
47         loss = loss_fcn(y_pred[train_mask], y_actual[train_mask])
48         optimizer.zero_grad()
49         loss.backward()
50         optimizer.step()
51         acc = evaluate(g, features, y_actual, val_mask, model)
52
53         accuracy_list.append(acc)
54         loss_list.append(loss.item())
55
56         print(
57             "Epoch {:05d} | Loss {:.4f} | Accuracy {:.4f} ".format(
58                 epoch, loss.item(), acc
59             )
60         )
61     return accuracy_list, loss_list
```

## Embedding generation algorithm (mini batch)

---

**Algorithm 2:** GraphSAGE minibatch forward propagation algorithm

---

**Input** : Graph  $\mathcal{G}(\mathcal{V}, \mathcal{E})$ ;  
input features  $\{\mathbf{x}_v, \forall v \in \mathcal{B}\}$ ;  
depth  $K$ ; weight matrices  $\mathbf{W}^k, \forall k \in \{1, \dots, K\}$ ;  
non-linearity  $\sigma$ ;  
differentiable aggregator functions  $\text{AGGREGATE}_k, \forall k \in \{1, \dots, K\}$ ;  
neighborhood sampling functions,  $\mathcal{N}_k : v \rightarrow 2^{\mathcal{V}}, \forall k \in \{1, \dots, K\}$   
**Output** : Vector representations  $\mathbf{z}_v$  for all  $v \in \mathcal{B}$

```
1  $\mathcal{B}^K \leftarrow \mathcal{B}$ ;  
2 for  $k = K \dots 1$  do  
3    $\mathcal{B}^{k-1} \leftarrow \mathcal{B}^k$ ;  
4   for  $u \in \mathcal{B}^k$  do  
5      $\mathcal{B}^{k-1} \leftarrow \mathcal{B}^{k-1} \cup \mathcal{N}_k(u)$ ;  
6   end  
7 end  
8  $\mathbf{h}_u^0 \leftarrow \mathbf{x}_v, \forall v \in \mathcal{B}^0$ ;  
9 for  $k = 1 \dots K$  do  
10  for  $u \in \mathcal{B}^k$  do  
11     $\mathbf{h}_{\mathcal{N}(u)}^k \leftarrow \text{AGGREGATE}_k(\{\mathbf{h}_{u'}^{k-1}, \forall u' \in \mathcal{N}_k(u)\})$ ;  
12     $\mathbf{h}_u^k \leftarrow \sigma(\mathbf{W}^k \cdot \text{CONCAT}(\mathbf{h}_u^{k-1}, \mathbf{h}_{\mathcal{N}(u)}^k))$ ;  
13     $\mathbf{h}_u^k \leftarrow \mathbf{h}_u^k / \|\mathbf{h}_u^k\|_2$ ;  
14  end  
15 end  
16  $\mathbf{z}_u \leftarrow \mathbf{h}_u^K, \forall u \in \mathcal{B}$ 
```

---

Make each mini batch  $\mathcal{B}^k, k \in \{1, \dots, K\}$

Identical to previous pseudocode

# Implement

29

## Define Neighborhood Sampler

```
1 from dgl.data.loading import DataLoader, NeighborSampler
2
3 sampler = NeighborSampler([5,8]) → Neighborhood sampler ( $S_1, S_2$ )
4 dataloader1 = DataLoader(
5     dataset1[0], torch.arange(dataset1[0].num_nodes()).to(dataset1[0].device), sampler, device=device,
6     batch_size=256, shuffle=False, drop_last=False, Node index, Sampler type
7     num_workers=0)
```

```
✓ [13] 1 for it, (input_nodes, output_nodes, blocks) in enumerate(dataloader1):
4空 2     print(blocks)

[Block(num_src_nodes=1779, num_dst_nodes=960, num_edges=4261), Block(num_src_nodes=960, num_dst_nodes=256, num_edges=1053)]
[Block(num_src_nodes=1814, num_dst_nodes=978, num_edges=4349), Block(num_src_nodes=978, num_dst_nodes=256, num_edges=1059)]
[Block(num_src_nodes=1801, num_dst_nodes=935, num_edges=4195), Block(num_src_nodes=935, num_dst_nodes=256, num_edges=996)]
[Block(num_src_nodes=1783, num_dst_nodes=882, num_edges=3922), Block(num_src_nodes=882, num_dst_nodes=256, num_edges=882)]
[Block(num_src_nodes=1758, num_dst_nodes=904, num_edges=4042), Block(num_src_nodes=904, num_dst_nodes=256, num_edges=986)]
[Block(num_src_nodes=1802, num_dst_nodes=951, num_edges=4164), Block(num_src_nodes=951, num_dst_nodes=256, num_edges=1041)]
[Block(num_src_nodes=1747, num_dst_nodes=933, num_edges=4218), Block(num_src_nodes=933, num_dst_nodes=256, num_edges=1182)]
[Block(num_src_nodes=1544, num_dst_nodes=841, num_edges=3753), Block(num_src_nodes=841, num_dst_nodes=256, num_edges=1209)]
[Block(num_src_nodes=1494, num_dst_nodes=797, num_edges=3616), Block(num_src_nodes=797, num_dst_nodes=256, num_edges=996)]
[Block(num_src_nodes=1419, num_dst_nodes=756, num_edges=3065), Block(num_src_nodes=756, num_dst_nodes=256, num_edges=783)]
[Block(num_src_nodes=632, num_dst_nodes=355, num_edges=1061), Block(num_src_nodes=355, num_dst_nodes=148, num_edges=301)]
```

# Implement

30

## Train (mini batch)

```
33 def train(model, dataloader):
34     # define train/val samples, loss function and optimizer
35     loss_fcn = nn.CrossEntropyLoss()
36
37     optimizer = torch.optim.Adam(model.parameters(), lr=1e-2, weight_decay=5e-4)
38
39     train_loss_list = []
40     val_acc_list = []
41
42     # training loop
43     for epoch in range(200):
44         model.train()
45         val_acc = 0
46         train_loss = 0
47
48         for input_nodes, output_nodes, batch_graphs in dataloader:
49
50             features = batch_graphs[0].srcdata['feat']
51             labels = batch_graphs[1].dstdata['label']
52             train_mask, val_mask = batch_graphs[1].dstdata["train_mask"], batch_graphs[1].dstdata["val_mask"]
53             #in_size = features.shape[1] = 1433 # number of features
54             #out_size = dataset1.num_classes = 7 # number of classes
55
56             logits = model(batch_graphs, features)
57             loss = loss_fcn(logits[train_mask], labels[train_mask])
58             optimizer.zero_grad()
59             loss.backward()
60             optimizer.step()
```


→ Batch training



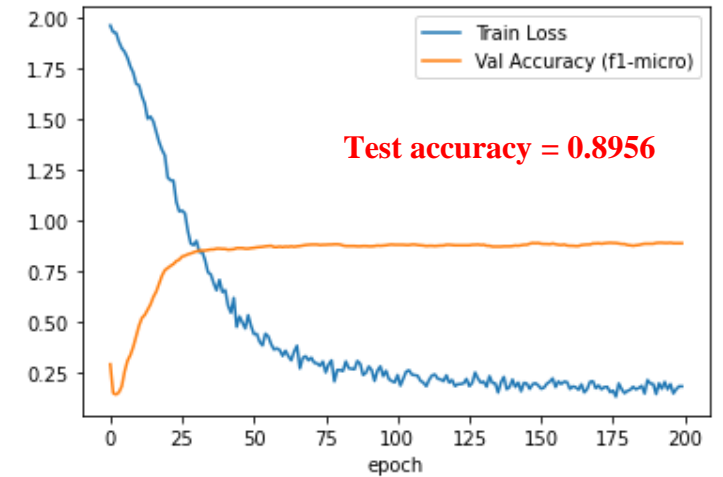
# Implement

31


## Cora Citation Dataset Result (full batch)



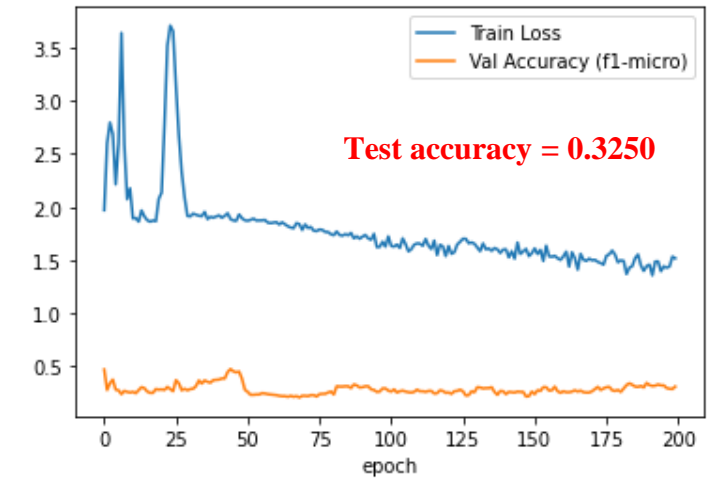
GraphSAGE-GCN



GraphSAGE-mean



GraphSAGE-max pooling




GraphSAGE-LSTM


# Implement

32


## Cora Citation Dataset Result (mini batch)




GraphSAGE-GCN



GraphSAGE-mean




GraphSAGE-max pooling



GraphSAGE-LSTM

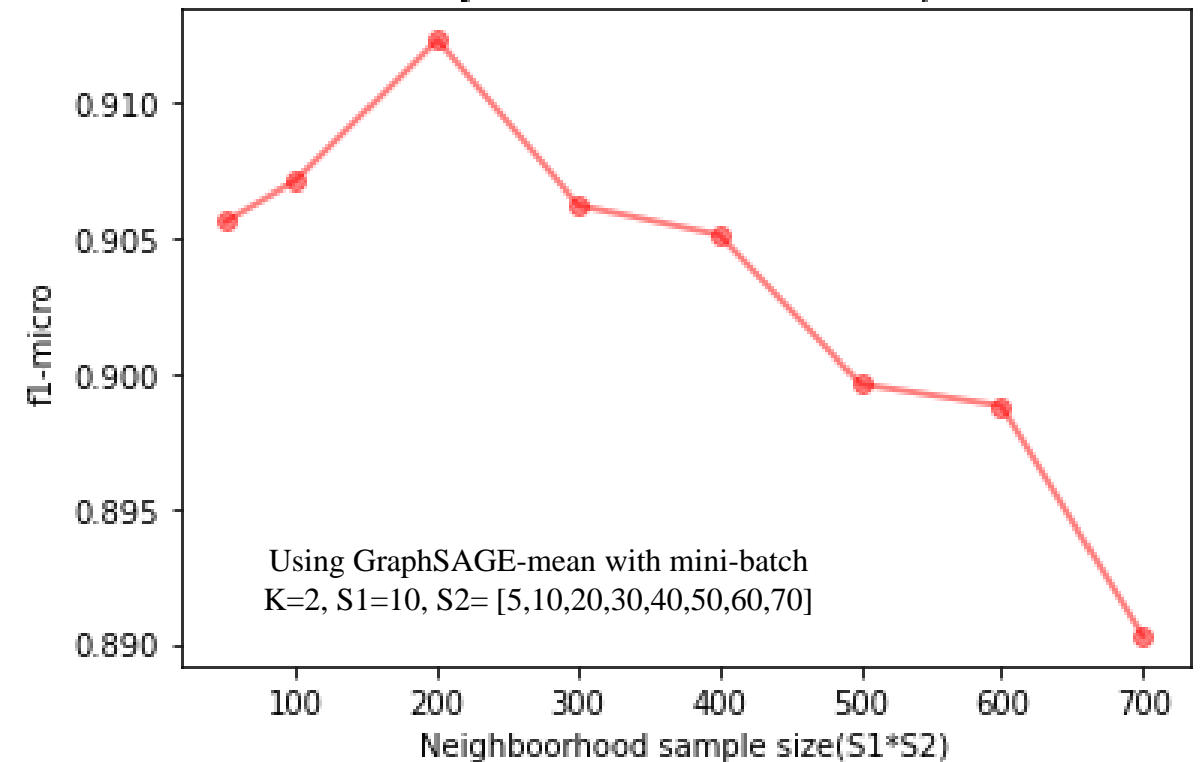
## Runtime Result Cora Citation Dataset

### Training Runtimes



## Accuracy over Neighborhood Sample size

### Accuracy over different sample size



Thank you