

Variational Autoencoder (VAE) & Graph Variational Autoencoder(GVAE)

February 7, 2023

Contents


1. Introduction
2. Structure of VAE
3. Structure of GVAE
4. Implementation
5. Conclusion

Introduction

Variational Inference + AutoEncoder = VAE

Variational Inference : Approximating target distribution with a simpler distribution


ex) 사후확률(posterior) 분포 $p(z|x) \rightarrow q(z) \sim \text{Normal}$



Introduction

Variational Inference + AutoEncoder = VAE

AutoEncoder : "The aim of an **autoencoder** is to learn a **representation** (encoding) for a set of data"

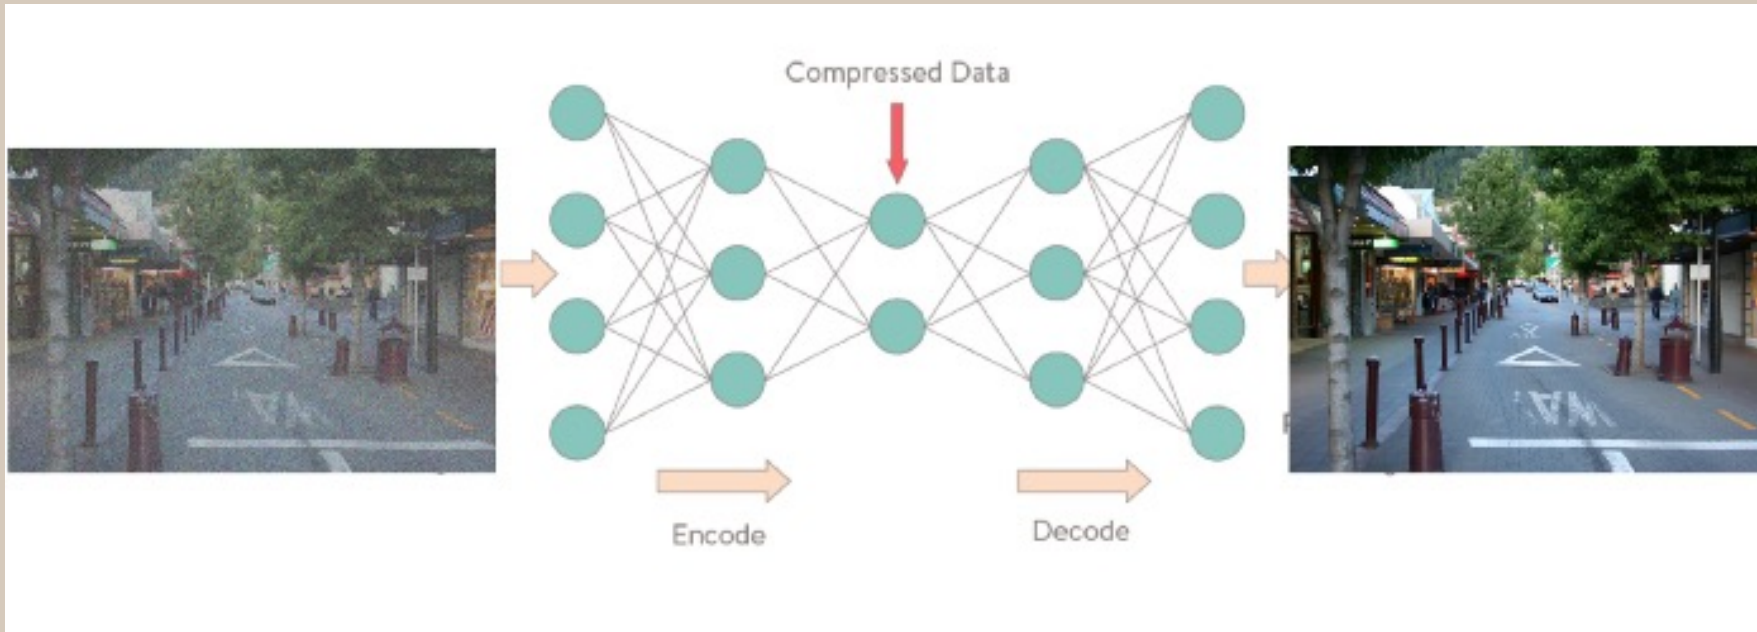


AutoEncoder (AE)

Introduction

Variational Inference + AutoEncoder = VAE

AutoEncoder : "The aim of an **autoencoder** is to learn a **representation (encoding)** for a set of data "





Common Applications:

1. Anomaly Detection
2. Image Denoising

Introduction

Why VAE?



Introduction

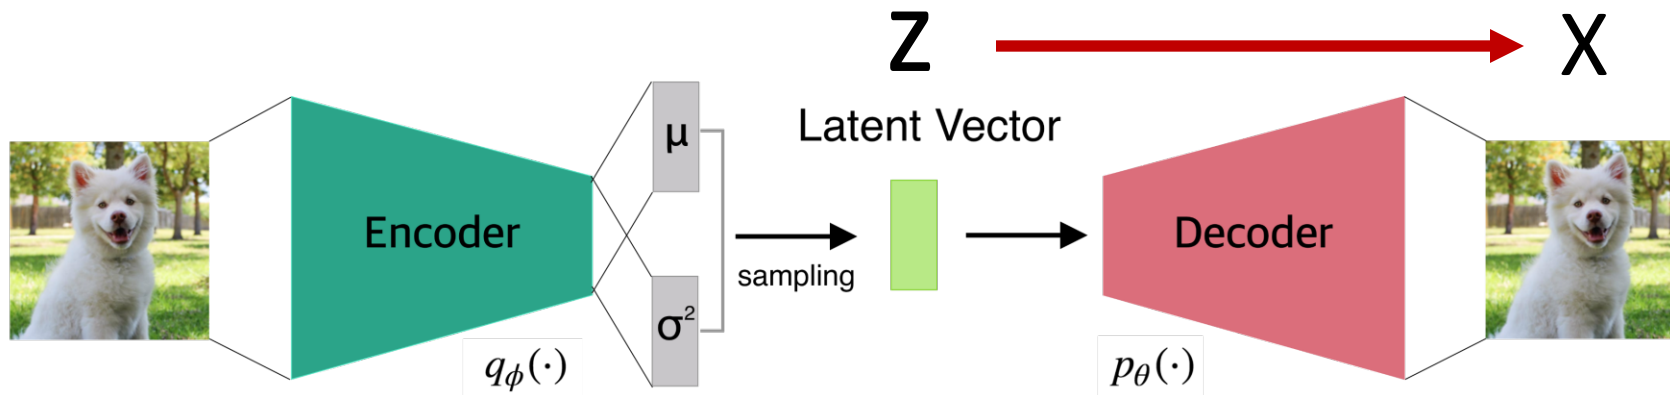
(1) Encoder

- inference network
- input : x , output : z
- $q(z \mid x, \phi)$


(2) Decoder

- generative network
- input : z , output : as closely as x
- $p(x \mid z, \theta)$

Variational Autoencoder



VAE



Given I.I.D data, we want to learn model parameters to maximize likelihood of training data

$$X_1, X_2, \dots, X_n$$



$$p_{\theta}(\mathbf{x}) = \int p_{\theta}(\mathbf{z})p_{\theta}(\mathbf{x}|\mathbf{z}) d\mathbf{z}$$

VAE

Data Likelihood : $p_{\theta}(\mathbf{x}) = \int p_{\theta}(\mathbf{z}) p_{\theta}(\mathbf{x}|\mathbf{z}) d\mathbf{z}$

Intractable to compute for every \mathbf{z}

Define additional encoder network $q(\mathbf{z}|\mathbf{x})$ that approximate $p(\mathbf{x}|\mathbf{z})$



VAE

$$\begin{aligned}\log p_{\theta}(\mathbf{x}) &= \int q_{\phi}(\mathbf{z}|\mathbf{x}) \log p_{\theta}(\mathbf{x}) d\mathbf{z} \\&= \int q_{\phi}(\mathbf{z}|\mathbf{x}) \log \frac{p_{\theta}(\mathbf{x}|\mathbf{z})p(\mathbf{z})}{p_{\theta}(\mathbf{z}|\mathbf{x})} d\mathbf{z} \\&= \int q_{\phi}(\mathbf{z}|\mathbf{x}) \log \frac{p_{\theta}(\mathbf{x}|\mathbf{z})p(\mathbf{z})}{p_{\theta}(\mathbf{z}|\mathbf{x})} \frac{q_{\phi}(\mathbf{z}|\mathbf{x})}{q_{\phi}(\mathbf{z}|\mathbf{x})} d\mathbf{z} \\&= \int q_{\phi}(\mathbf{z}|\mathbf{x}) \log p_{\theta}(\mathbf{x}|\mathbf{z}) d\mathbf{z} - KL(q_{\phi}(\mathbf{z}|\mathbf{x})||p(\mathbf{z})) + KL(q_{\phi}(\mathbf{z}|\mathbf{x})||p_{\theta}(\mathbf{z}|\mathbf{x}))\end{aligned}$$

$$D_{KL}(P||Q) = \int P(x) \log \frac{P(x)}{Q(x)} dx$$

- $KL(q||p) \geq 0$
- $KL(q||p) = 0 \Leftrightarrow q = p$

VAE

$$= \int q_{\phi}(\mathbf{z}|\mathbf{x}) \log p_{\theta}(\mathbf{x}|\mathbf{z}) d\mathbf{z} - KL(q_{\phi}(\mathbf{z}|\mathbf{x})||p(\mathbf{z})) + \underbrace{KL(q_{\phi}(\mathbf{z}|\mathbf{x})||p_{\theta}(\mathbf{z}|\mathbf{x}))}_{\text{(Non-Negative)}}$$

$$\geq \int q_{\phi}(\mathbf{z}|\mathbf{x}) \log p_{\theta}(\mathbf{x}|\mathbf{z}) d\mathbf{z} - KL(q_{\phi}(\mathbf{z}|\mathbf{x})||p(\mathbf{z})) \leftarrow \text{Evidence lower bound (ELBO)}$$

Reconstruction error **Regularization error**



**Maximizing Evidence
Lower Bound(ELBO)**

probability a hypothesis is true
given the evidence

probability a hypothesis is true
(before any evidence is present)

probability of seeing the evidence
if the hypothesis is true

probability of observing the evidence

$$P(\textcolor{red}{H}/\textcolor{green}{E}) = \frac{P(\textcolor{red}{H}) P(\textcolor{green}{E}/\textcolor{red}{H})}{P(\textcolor{green}{E})}$$

VAE



– ELBO


$$\arg \min_{\theta, \phi} \sum_i \left[-\mathbb{E}_{q_{\phi}(z|x_i)} [\log(p(x_i|g_{\theta}(z)))] + KL(q_{\phi}(z|x_i)||p(z)) \right]$$

Reconstruction Error

Regularization

VAE

$$\arg \min_{\theta, \phi} \sum_i \underbrace{-\mathbb{E}_{q_{\phi}(z|x_i)} [\log(p(x_i|g_{\theta}(z)))]}_{\text{Evidence Lower Bound (ELBO)}} + KL(q_{\phi}(z|x_i)||p(z))$$



$$\mathbb{E}_{q_{\phi}(z|x_i)} [\log(p_{\theta}(x_i|z))] = \int \log(p_{\theta}(x_i|z)) q_{\phi}(z|x_i) dz \approx \frac{1}{L} \sum_{z^{i,l}} \log(p_{\theta}(x_i|z^{i,l}))$$

VAE

$$\mathbb{E}_{q_{\phi}(z|x_i)}[\log(p_{\theta}(x_i|z))] = \int \log(p_{\theta}(x_i|z))q_{\phi}(z|x_i)dz \approx \frac{1}{L} \sum_{z^{i,l}} \log(p_{\theta}(x_i|z^{i,l})) \approx \log(p_{\theta}(x_i|z^i))$$

$$p_{\theta}(x_i|z^i) \sim \text{Bernoulli}(p_i)$$

$$= \sum_{j=1}^D \log p_{i,j}^{x_{i,j}} (1 - p_{i,j})^{1-x_{i,j}} \quad \leftarrow p_{i,j}: \text{network output}$$

$$= \sum_{j=1}^D x_{i,j} \log p_{i,j} + (1 - x_{i,j}) \log(1 - p_{i,j})$$

Cross entropy

VAE

$$\log(p_{\theta}(x_i|z^i)) = \log(N(x_i; \mu_i, \sigma_i^2 I))$$

$$= - \sum_{j=1}^D \frac{1}{2} \log(\sigma_{i,j}^2) + \frac{(x_{i,j} - \mu_{i,j})^2}{2\sigma_{i,j}^2}$$

For gaussian distribution with identity covariance

$$\log(p_{\theta}(x_i|z^i)) \propto - \sum_{j=1}^D (x_{i,j} - \mu_{i,j})^2 \quad \text{Squared Error}$$

VAE

Optimization Assumptions

$$q_{\phi}(z|x_i) \sim N(\mu_i, \sigma_i^2 I)$$

$$p(z) \sim N(0, I)$$

$$\arg \min_{\theta, \phi} \sum_i -\mathbb{E}_{q_{\phi}(z|x_i)} [\log(p(x_i|g_{\theta}(z)))] + \boxed{KL(q_{\phi}(z|x_i) || p(z))}$$

Regularization

Theorem: Let x be an $n \times 1$ random vector. Assume two multivariate normal distributions P and Q specifying the probability distribution of x as

$$\begin{aligned} P : x &\sim \mathcal{N}(\mu_1, \Sigma_1) \\ Q : x &\sim \mathcal{N}(\mu_2, \Sigma_2) . \end{aligned} \tag{1}$$

Then, the Kullback-Leibler divergence of P from Q is given by


$$\text{KL}[P || Q] = \frac{1}{2} \left[(\mu_2 - \mu_1)^T \Sigma_2^{-1} (\mu_2 - \mu_1) + \text{tr}(\Sigma_2^{-1} \Sigma_1) - \ln \frac{|\Sigma_1|}{|\Sigma_2|} - n \right] . \tag{2}$$

VAE


$$\begin{aligned} KL(q_{\phi}(z|x_i)||p(z)) &= \frac{1}{2} \left\{ \text{tr}(\sigma_i^2 I) + \mu_i^T \mu_i - J + \ln \frac{1}{\prod_{j=1}^J \sigma_{i,j}^2} \right\} \\ &= \frac{1}{2} \left\{ \sum_{j=1}^J \sigma_{i,j}^2 + \sum_{j=1}^J \mu_{i,j}^2 - J - \sum_{j=1}^J \ln(\sigma_{i,j}^2) \right\} \\ &= \frac{1}{2} \sum_{j=1}^J (\mu_{i,j}^2 + \sigma_{i,j}^2 - \ln(\sigma_{i,j}^2) - 1) \end{aligned}$$

VAE

VAE Structure



VAE




Backpropagation

VAE


$$L_i(\phi, \theta, x_i) = \underbrace{-\mathbb{E}_{q_\phi(z|x_i)} [\log(p_\theta(x_i|z))] }_{\text{Reconstruction Error}} + \underbrace{KL(q_\phi(z|x_i)||p(z))}_{\text{Regularization}}$$

Reconstruction Error

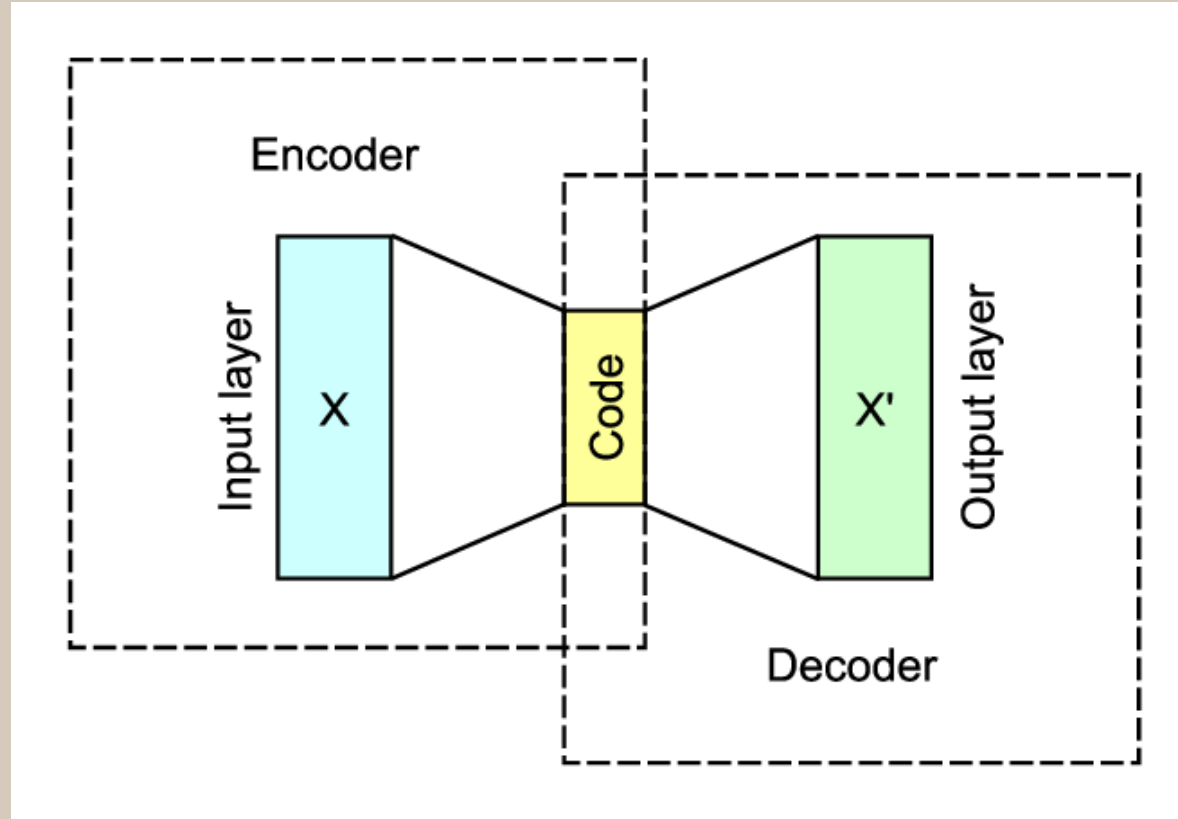
Regularization


$$\log(p_\theta(x_i|z^i)) = \sum_{j=1}^D x_{i,j} \log p_{i,j} + (1 - x_{i,j}) \log(1 - p_{i,j})$$

Cross entropy


$$KL(q_\phi(z|x_i)||p(z)) = \frac{1}{2} \sum_{j=1}^J (\mu_{i,j}^2 + \sigma_{i,j}^2 - \ln(\sigma_{i,j}^2) - 1)$$

GVAE



GVAE

First Neural Net of Encoder

$$\tilde{A} = D^{-1/2} A D^{-1/2}$$

$$\bar{X} = GCN(A, X) = ReLU(\tilde{A} X W_0)$$

GVAE

First Neural Net of Encoder


$$\tilde{A} = D^{-1/2} A D^{-1/2}$$

$$\bar{X} = GCN(A, X) = ReLU(\tilde{A} X W_0)$$

Second Neural Net of Encoder

$$\mu = GCN_{\mu}(X, A) = \tilde{A} \bar{X} W_1$$


$$\log \sigma^2 = GCN_{\sigma}(X, A) = \tilde{A} \bar{X} W_1$$



$$GCN(A, X) = \tilde{A} ReLU(\tilde{A} X W_0) W_1$$

with $\tilde{A} = D^{-1/2} A D^{-1/2}$

GVAE



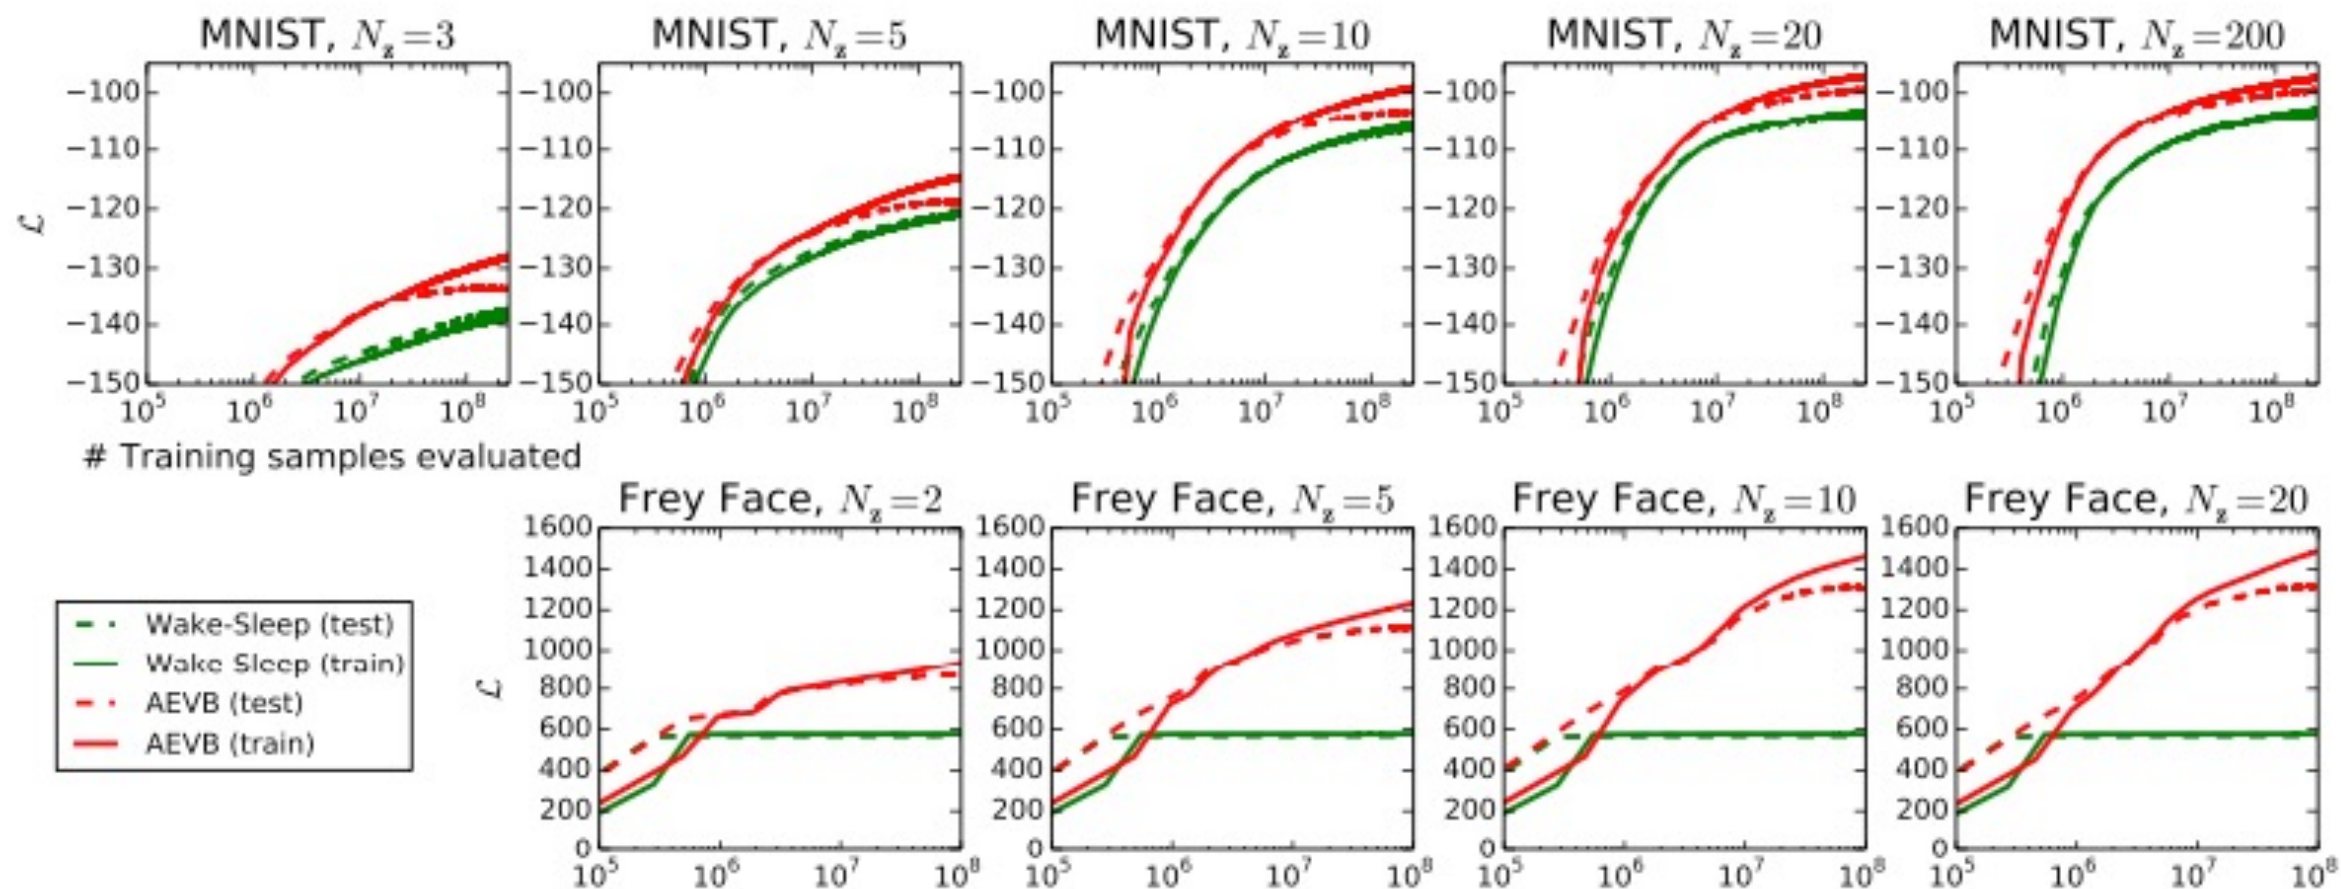
Inner product between
latent variable Z

$$\hat{A} = \text{logistic sigmoid}(zz^T)$$


$$GCN(A, X) = \tilde{A} \text{ReLU}(\tilde{A} X W_0) W_1$$

$$Z = \mu + \sigma \odot \epsilon$$
$$\epsilon \sim \text{Norm}(0, 1)$$

Experiment




Experiment



Implementation – VAE

```
45 for i in train_loader:
46     plt.figure(figsize=(10,10))
47     k = None
48     for img in i[0][5:10]:
49         img = img.permute(1, 2, 0)
50         img = img.squeeze()
51         if k == None:
52             k = img
53             print(k.shape)
54         else:
55             k = torch.cat((k, img), 1)
56     plt.imshow(k)
57     break
58 plt.show()
59
```



Implementation – VAE

```
60 class VAE(nn.Module):
61     def __init__(self, input_dim, hid_dim1, hid_dim2, emb_dim):
62         super(VAE, self).__init__()
63
64         self.input_dim = input_dim
65
66         # encoder part
67         self.fc1 = nn.Linear(input_dim, hid_dim1)
68         self.fc2 = nn.Linear(hid_dim1, hid_dim2)
69         self.fc3_mean = nn.Linear(hid_dim2, emb_dim)
70         self.fc3_var = nn.Linear(hid_dim2, emb_dim)
71         # decoder part
72         self.fc4 = nn.Linear(emb_dim, hid_dim2)
73         self.fc5 = nn.Linear(hid_dim2, hid_dim1)
74         self.fc6 = nn.Linear(hid_dim1, input_dim)
75
```

Implementation – VAE


```
76     def forward(self, x):
77
78         #encoding
79         temp = F.relu(self.fc1(x.view(-1, 784)))
80         temp = F.relu(self.fc2(temp))
81         mu = self.fc3_mean(temp)
82         log_var = self.fc3_var(temp)
83
84         #sampling
85         std = torch.exp(0.5*log_var) # This is to constrain the variance to be positive ( $\sigma^2 \in \mathbb{R}^+$ )
86                                     # But why learn log_var instead of learning log_sigma,
87                                     # which does not require 0.5 multiplication?
88         eps = torch.randn_like(std)
89         z = eps.mul(std).add_(mu)
90
91         #decoding
92         temp = F.relu(self.fc4(z))
93         temp = F.relu(self.fc5(temp))
94         return F.sigmoid(self.fc6(temp)), mu, log_var
95
```

Implementation – VAE

```
96 # build model
97 vae = VAE(784, 128, 32, 20)
98 if torch.cuda.is_available():
99     vae.cuda()
100
101 optimizer = optim.Adam(vae.parameters())
102
```

```
103 # return reconstruction error + KL divergence losses
104 def loss_function(reconstructed_data, original_data, mu, log_var):
105     Binary_CE = F.binary_cross_entropy(reconstructed_data, original_data.view(-1, 784), reduction='sum')
106     KL_Divergence = -0.5 * torch.sum(1 + log_var - mu**2 - log_var.exp())
107     return Binary_CE + KL_Divergence
```


Implementation – VAE



Implementation – GVAE

```
1 import torch
2 import torch.nn as nn
3 import torch.nn.functional as F
4 import numpy as np
5
6 class VGAE(nn.Module):
7     def __init__(self, adj):
8         super(VGAE, self).__init__()
9         self.first_layer = GraphConv(input_dim, output_dim*2, adj, F.relu)
10        self.mean = GraphConv(output_dim*2, output_dim, adj, activation=lambda x:x)
11        self.logvar = GraphConv(output_dim*2, output_dim, adj, activation=lambda x:x)
12
13
14    def forward(self, X):
15        hidden = self.first_layer(X)
16        self.mean = self.mean(hidden)
17        self.logvar = self.logvar(hidden)
18        epsilon = torch.randn(self.mean.size(0), output_dim)
19        sampled_z = epsilon*torch.exp(self.logvar) + self.mean
20        decoded = torch.sigmoid(torch.matmul(Z,Z.t()))
21        return decoded
22
```

Implementation – GVAE

```
23 class GraphConv(nn.Module):
24     def __init__(self, input_dim, output_dim, adj, activation, **kwargs):
25         super(GraphConv, self).__init__(**kwargs)
26         self.weight = torch.rand(input_dim, output_dim)
27         self.adj = adj
28         self.activation = activation
29
30     def forward(self, inputs):
31         X = inputs
32         XW = torch.mm(X, self.weight)
33         AXW = torch.mm(self.adj, XW)
34         output = self.activation(AXW)
35         return output
```

Conclusion

- VAE is a novel estimator of the variational lower bound, Stochastic Gradient VB (SGVB), for efficient approximate inference with continuous latent variables.
- Both GVAE and GAE achieve competitive results on the featureless task.
- Future work will investigate better-suited prior distributions, more flexible generative models and the application of a stochastic gradient descent algorithm for improved scalability.